staying ahead

ISSUE 003
MON, 27 APRIL 2026

TODAY'S HEADLINE

China just made the world's biggest AI free.

DeepSeek just dropped a free ChatGPT rival. One AI data center cost $16 billion. AI bots traded $4,000 with each other. Europe got its first $20B AI giant.

It has been a long 24 hours. Here is the rundown, in 7 minutes.

SENT BY
Vaibhav Sisinty

MORNING DIGEST

What we cover today.

Seven stories from the last 24 hours, plus a personal milestone. Read in this order and you will be ahead of 90% of people in AI by lunch.

01 **DeepSeek dropped a free ChatGPT rival**
1.6 trillion parameters. Free for anyone to download. Costs one-seventh of GPT 5.5.

02 **An OpenAI data center just cost $16 billion**
Related Digital and Blackstone. Built for Oracle. Powers OpenAI's Stargate.

03 **AI bots just traded $4,000 with each other**
186 deals. 69 employees. Zero humans in the loop. The losers never noticed.

04 **Europe just built its own $20B AI giant**
Cohere acquired Aleph Alpha. The first credible non-US AI counterweight at scale.

05 **Maine just killed America's first AI ban**
Governor vetoed the data center moratorium. AI build-out path stays open.

06 **Microsoft just made voice AI free**
VibeVoice covers TTS, ASR, and real-time streaming. Runs on a laptop.

07 **A new chip just cut AI energy use by 70%**
Cambridge lab. May flip the entire AI compute equation if it scales.

01 . FEATURE . FRONTIER MODEL

DeepSeek just dropped a free ChatGPT rival.

QUICK TAKE

DeepSeek dropped V4-Pro and V4-Flash on April 24. Both ship under MIT license with a 1 million token context window. V4-Pro is 1.6 trillion parameters total, the largest open weights model in existence. Pricing comes in at one-seventh of GPT 5.5.

WHAT YOU NEED TO KNOW

- **The architecture:** 1.6T total params, 49B active per token. V4-Flash is 284B / 13B active. Both run 1M context with dual Thinking and Non-Thinking modes.
- **The pricing:** Flash at $0.14 in / $0.28 out per million tokens. Pro at $1.74 / $3.48. The cheapest frontier model on the market.
- **Migration:** keep your base_url, change the model parameter to deepseek-v4-pro or deepseek-v4-flash. Old aliases retire 24 July 2026.
- **The benchmark:** V4-Pro hits a Codeforces rating of 3,206. That ranks 23rd among human competitors on the live leaderboard.

THE WILD PART

V4-Flash weights are 160GB on Hugging Face. **An indie team can self-host frontier intelligence on a multi-GPU mid-size box this weekend.** The cost wall just collapsed.

02 . INVESTMENT . AI INFRASTRUCTURE

An OpenAI data center just cost $16 billion.

QUICK TAKE

Related Digital and Blackstone closed $16B in financing on April 24 for a single data center in Saline Township, Michigan. Purpose-built for Oracle, leased primarily to support OpenAI workloads under Stargate. More than 1 gigawatt of capacity. One of the largest single-project data center financings in US history.

WHAT YOU NEED TO KNOW

- **The deal:** $16B total. ~$14B in fixed-rate, long-term debt anchored by PIMCO. Equity from Related Digital and Blackstone. Bank of America structured the deal.
- **The capacity:** 1+ gigawatt. About 25% of DTE Energy's current peak capacity. Typical data centers are 100 to 300 MW. This is hyperscale.
- **The chain:** Related Digital builds it. Oracle leases it. OpenAI runs on it. Part of the Stargate compute expansion.
- **The pattern:** This follows Oracle's $38B Texas and Wisconsin debt deal and an $18B New Mexico site. Compute is now an asset class.

THE WILD PART

$16B for one site is what entire chip foundries used to cost. Compute capacity is the new steel mill. **The capital floor for staying in the AI game just moved again.**

03 . NEWS . AGENT ECONOMY

AI bots just traded $4,000 with each other.

READ ON X . @VAIBHAVSISINTY →

QUICK TAKE

Anthropic ran an experiment called Project Deal. 69 employees, $100 each in gift cards, Claude agents handling all negotiations. The agents closed 186 deals across 500+ items. Then Anthropic ran a parallel test that quietly exposed a problem.

WHAT YOU NEED TO KNOW

- **The setup:** a Slack-based classified marketplace. Claude interviewed each employee to capture preferences, then agents listed, negotiated, and closed deals autonomously.
- **The volume:** 186 completed deals, $4,000+ in real money moved. 46% of participants said they would pay for the service.
- **The hidden test:** half of users got Opus 4.5, half got Haiku 4.5. Opus sellers earned $2.68 more per item. Opus buyers saved $2.45.
- **The fairness gap:** Haiku users rated their deals just as fair as Opus users. They never noticed they were losing.

THE WILD PART

When agents trade with agents, model quality decides who wins. **The losers do not feel the loss.** That is a new kind of inequality, invisible by design.

04 . FUND RAISING . SOVEREIGN AI

Europe just built its own $20B AI giant.

QUICK TAKE

Cohere agreed to acquire Germany's Aleph Alpha on April 24. The combined entity is valued at $20B, anchored in Canada and Germany, blessed by both governments. Schwarz Group, the owner of Lidl, leads Cohere's upcoming Series E with $600M.

WHAT YOU NEED TO KNOW

- **The structure:** Cohere shareholders take ~90% of the merged entity. Aleph Alpha takes ~10%. Dual headquarters in Canada and Germany.
- **The capital:** Schwarz Group commits $600M (~€500M) as lead investor in Cohere's Series E, due to close in 2026.
- **The customers:** Aleph Alpha brings Deutsche Bank, Bosch, SAP, German government. Cohere brings Royal Bank of Canada, Salesforce, Oracle.
- **The trade:** Schwarz expects the new entity to run on STACKIT, its sovereign cloud platform. A retail giant just bought itself a captive AI customer.

THE WILD PART

This is the first credible non-US AI counterweight at scale. **If you sell into regulated EU enterprise, your deck just got a new logo to defend against.**

05 . NEWS . POLICY

Maine just killed America's first AI ban.

QUICK TAKE

Gov. Janet Mills vetoed L.D. 307 on April 24. The bill would have paused permits for new data centers above 20 MW until November 2027, the first such moratorium in any US state. It cleared both chambers with bipartisan support before her veto.

WHAT YOU NEED TO KNOW

- **The bill:** 18-month pause on data centers above 20 MW. Would have been the first state-level AI compute moratorium in the country.
- **The reason:** Mills wanted a carve-out for a $550M data center project at the former Androscoggin Mill in Jay. The bill did not include it.
- **The next move:** Mills will create a 13-member data center council via executive order to study impacts on grid, water, and ratepayers.
- **The watch:** New York has a similar three-year moratorium proposal. With Maine off the board, that one is now the next test.

THE WILD PART

AI compute is the new pipeline politics. State legislatures now treat data centers the way they once treated refineries. **The fight has moved from model floors to power grids.**

06 . OPEN SOURCE . BUILDER'S TOOLKIT

Microsoft just made voice AI free.

QUICK TAKE

Microsoft's VibeVoice family quietly hit critical mass this month. It now covers TTS, ASR, and a real-time streaming variant. ASR handles 60-minute audio in a single pass with speaker diarization. All MIT licensed, all on GitHub.

WHAT YOU NEED TO KNOW

- **The TTS:** synthesizes 90-minute multi-speaker dialogue with up to 4 distinct voices. Optimized for podcast and audiobook generation.
- **The ASR:** 60-minute single-pass transcription with Who, When, What output (speaker, timestamp, content). 50+ languages supported.
- **The realtime variant:** VibeVoice-Realtime-0.5B streams TTS at low latency. Suitable for voice agents and live chatbots.
- **The license:** MIT. Runs offline on consumer hardware. 4GB GPU RAM is enough.

THE WILD PART

Voice AI just stopped being an API line item. If your stack pays ElevenLabs or Whisper today, you can move to a self-hosted Microsoft model this week. **The voice agent moat is now a commodity.**

07 . RESEARCH PAPER . FROM THE LABS

A new chip just cut AI energy use by 70%.

QUICK TAKE

A team at Cambridge published a new memristor design in Science Advances on April 23. The chip uses a hafnium oxide thin film to switch states at currents about a million times lower than conventional memristors. The projected energy reduction for AI inference is up to 70%.

WHAT YOU NEED TO KNOW

- **The science:** hafnium oxide doped with strontium and titanium, grown in two stages. Creates p-n junctions that switch via energy barriers, not unstable filaments.
- **The advantage:** switching currents about a million times lower than current oxide-based memristors. Outstanding device-to-device uniformity.
- **The catch:** endurance is still 10,000 to 100,000 cycles. Retention is also limited. Not yet ready for production-scale storage.
- **The IP:** patent already filed via Cambridge Enterprise. The team led by Dr. Babak Bakhit spent three years on the breakthrough.

THE WILD PART

Inference cost is the next AI ceiling. If neuromorphic chips like this scale, the entire compute equation flips. **GPU shortages would matter less than fab capacity for new substrates.**

EDITORIAL . MY TWO CENTS

One pattern. Told seven different ways.

The last 24 hours were not a list of unrelated stories. Read in sequence, they are one shift, told in seven voices.

DeepSeek gave the world a free ChatGPT rival. $16 billion got committed to one OpenAI data center. AI bots traded real money with other AI bots. Europe built its first $20 billion AI giant. Maine killed America's first AI ban. Microsoft made voice AI free. A new chip cut AI energy use by 70%.

> **This is not seven stories. This is the same story told seven times. The cost of intelligence is collapsing, the locus of control is shifting, and the moat is moving from models to compute, sovereignty, and energy.**

The labs that win the next two years will not be the ones with the best benchmark. They will be the ones that locked compute, locked sovereignty deals, and locked energy. Everyone else is renting.

IF YOU TAKE ONE THING FROM THIS

Stop building on a single model. Build the workflow that lets you swap any model in. The cost wall is collapsing. Your moat is the system around the model, not the model itself.

Until tomorrow,

Vaibhav

MILESTONE . FROM THE FOUNDER

500K subscribers. In 6 months.

AIA UPDATES ON YOUTUBE

6 months ago, we decided to start YouTube to teach AI in a deeper way. Instagram was great, but 3-minute reels were not enough.

The blocker was my time. I travel constantly. Running a 200 person team is already a full-time job.

So we built an AI avatar of me. The team syncs with me 2 to 3 hours a week. My thoughts and ideas get baked into the content. I do not physically shoot a single video.

The real question was: would anyone watch 20-minute long-form videos built around an AI avatar?

Today, we crossed 500K. Probably the only YouTube channel in the world at 500K where 100% of the content runs on the founder's AI avatar.

Not everyone loves the avatar part. **But millions of people are learning AI through it. That is what matters.**

Next goal: 1M subs. Then the biggest AI education channel in the world. We are just getting started.

100%	0	1M
AI avatar generated content	Videos physically shot	Next subscriber milestone

FROM THE FOUNDER

Vaibhav

THE WEEKLY DIGEST

Catch up on the last 7 days.

Everything that hit the channel this week. Tap any thumbnail to open it on YouTube.

Youtube

This week's drops

@vaibhavsisinty →

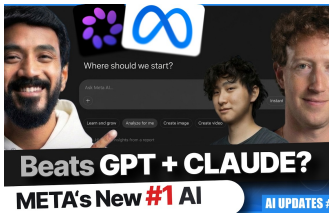

YOUTUBE Watch →

YOUTUBE Watch →

YOUTUBE Watch →

YOUTUBE Watch →

Youtube

Also worth your time

@vaibhavsisinty →

YOUTUBE Watch →

YOUTUBE Watch →

YOUTUBE Watch →

EVERY MORNING . 10:00 AM IST

You are caught up. Now *stay* caught up.

This drops in the **Outskill AI community** every morning at **10 AM IST**, before it goes anywhere else. Ten minutes. Zero fluff. The exact stack you need to stay one day ahead of everyone else in AI.

Primary . Free Forever

Join the WhatsApp community.

Tomorrow's PDF, the day after's, and every drop after that. In your pocket before anyone on LinkedIn sees it.

join.switchit.app/YT1

→

Watch

AIA Updates on YouTube

youtube.com/@vaibhavsisinty →

Get Featured

Want your tool in tomorrow's digest?

hi@growthschool.io →

↯ **Forward this PDF** to one builder, founder, or operator who should also be reading it. **That is how StayingAhead grows.**